

Yasho Ramith



9566054524



yashoramith@gmail.com



linkedin.com/in/yasho-ramith-a662622ab

Education

Birla Institute of Technology and Science, Pilani

Goa Campus

Bachelor of Technology

Computer Science and Engineering

2017 – 2021

Projects

Inference Engineering Platform Python / FastAPI / PyTorch / RAG

- Designed a production-oriented **LLM inference layer** focused on reducing serving overhead through model routing, request batching, response caching, concurrency controls, and token-aware execution.
- Built asynchronous inference paths separating request acceptance from model execution, allowing concurrent workloads to scale independently while supporting streaming responses and controlled backpressure.
- Optimized retrieval-to-generation execution using **vector retrieval, context filtering, caching, and output validation**, reducing unnecessary model work while maintaining grounded responses.
- Added execution telemetry covering **latency, token consumption, model usage, cache behavior, throughput, and failure modes** to analyze performance and cost characteristics of model-backed workloads.

Harness Bench Python / MCP / LLM Evaluation / Tracing

- Built an **agent evaluation harness** that replays production-like tasks and evaluates complete execution traces rather than scoring only final LLM responses.
- Implemented evaluation across **tool selection, retrieval quality, groundedness, task completion, response quality, latency, and failure modes** for MCP-enabled agent workflows.
- Converted individual agent failures into repeatable regression cases for prompts, tools, retrieval configuration, and execution logic through structured scorecards and trace analysis.
- Improved evaluation quality from approximately **0.79 to 0.93** through iterative evaluation, failure analysis, and targeted changes to prompts, tool definitions, and agent context.

Jevis AI – Enterprise Agent Platform FastAPI / RAG / MCP / LLMs

- Built an **enterprise AI agent platform** that converts natural-language requests into multi-step operational workflows using LLM reasoning, retrieval, MCP tools, and structured backend actions.
- Designed execution around **retrieval, planning, tool selection, validation, permissions, and human approval**, keeping agent behavior decoupled from underlying service implementations.
- Implemented **RAG and vector retrieval** to ground agent decisions in application data and assemble task-specific context before invoking downstream tools.
- Added **streaming execution, audit trails, tool permissions, failure recovery, and evaluation hooks**, making multi-step agent behavior observable and controllable.

Technical Skills

Languages

Python, Java, TypeScript, JavaScript, SQL, C++

AI / LLM

LLMs, RAG, AI Agents, Agentic AI, Claude, Prompt Engineering, Context Engineering, Structured Outputs

Agent Systems

MCP, Tool Calling, Function Calling, Planning, Memory, Guardrails, Human-in-the-Loop

Inference / Evaluation

Inference Optimization, Model Routing, Token Optimization, LLM Evaluation, Groundedness, Retrieval Evaluation

Backend / Distributed Systems

Java, Spring Boot, Spring Cloud, FastAPI, Node.js, REST, GraphQL, Microservices, Kafka, Asynchronous Processing

Data Platforms

Apache Spark, Databricks, BigQuery, SQL Optimization, ETL, Data Pipelines, Data Migration, Data Quality

Experience

Salesforce

Aug 2024 – Present

Senior Member of Technical Staff (SMTS) – AI Platform Engineering
Bangalore, India

- Reduced backend latency by approximately 35%** for the enterprise AI execution platform by optimizing model-backed workflows through model routing, request caching, concurrency controls, and token-aware execution.
- Increased processing throughput approximately 3×** for asynchronous AI operations by separating long-running agent execution from synchronous API traffic and coordinating work through Kafka, Redis, GraphQL, and asynchronous worker services.
- Designed the platform's **agent execution pipeline** around planning, retrieval, tool selection, validation, retries, and approval, turning natural-language requests into controlled operations against enterprise backend capabilities.
- Built reusable **MCP interfaces** over enterprise APIs and data sources, giving agents standardized contracts for discovering and invoking business operations without coupling agent logic to individual service implementations.
- Built the **evaluation and observability layer** for these workflows, measuring retrieval relevance, groundedness, task completion, tool-selection behavior, latency, token consumption, and failure categories across agent executions.
- Connected streaming AI interfaces to the same asynchronous execution pipeline, exposing task state and execution progress while backend workers continued processing long-running operations.
- Productionized the platform using **Docker, Kubernetes, GitHub Actions, and ArgoCD**, adding unit, integration, API, and AI-evaluation gates to the deployment path.

Salesforce

Jun 2023 – Aug 2024

Member of Technical Staff (MTS) – Backend & Distributed Systems
Bangalore, India

- Owned backend processing for an **enterprise event-driven workflow**, where synchronous service requests initiated downstream processing through Kafka; redesigned the ingestion-to-processing boundary and **improved throughput approximately 60%**.
- Separated **event ingestion, downstream processing, and failure recovery** so processing stages could scale independently and transient downstream failures would not block upstream service traffic.
- Implemented the service layer with **Java 17, Spring Boot, and Spring Cloud**, connecting REST APIs to event contracts, asynchronous workers, processing state, and downstream enterprise services.
- Reduced AWS Lambda cold-start latency approximately 40%** by restructuring initialization and integration boundaries within frequently invoked backend execution paths.
- Built validation for **AI-enabled backend workflows**, using prompt variants, multi-turn scenarios, response scoring, regression cases, and reliability checks to detect behavior changes before production releases.
- Instrumented the same request-to-processing path with **OpenTelemetry and CloudWatch**, correlating service telemetry and downstream dependencies to isolate latency buildup and failure boundaries.
- Owned changes across **API design, event-flow design, implementation, automated testing, deployment, and production troubleshooting**, taking backend workflows from requirements to operational production systems.

Walmart

Jun 2021 – Jun 2023

Software Engineer – Data Platform, Fraud Analytics & Backend India

- Engineered distributed processing for **transaction and fraud analytics workloads**, using Spark-based transformations, partitioning, and query optimization across datasets containing **hundreds of millions of records**.
- For telemetry retention and cloud-cost optimization, designed an archival pipeline handling **69.8M+ AppTraces records across 180 days**, moving historical telemetry from Log Analytics into Blob Storage.
- Built **50K–100K record extraction batches** with offset tracking, checkpointing, deterministic sequencing, validation, and retry handling, making large exports restartable instead of restarting the complete workload after a failed batch.
- Designed the archival workflow for **~1 TB-scale data movement**, separating query execution, batch persistence, validation, and reten-

Cloud / Infrastructure

AWS, Azure, GCP, Docker, Kubernetes, GitHub Actions, ArgoCD, CI/CD

Storage / Services

Azure Blob Storage, Log Analytics, Cosmos DB, Redis, PostgreSQL, MongoDB, DynamoDB

Observability

KQL, OpenTelemetry, CloudWatch, Distributed Tracing, Production Monitoring

tion to isolate failures and maintain incremental progress.

- Rebuilt the **fraud-reporting data path** around T360 transaction datasets, replacing legacy ECOMM sources across orders, order lines, hold types, and release states and feeding the migrated data into downstream Tableau reporting.
- Converted existing distributed-processing logic into **BigQuery-native transformations**, validating source-to-target relationships and downstream reporting behavior before switching the reporting workflow to the migrated source.
- Worked across **data processing, cloud storage, service integration, access control, and CI/CD**, resolving production issues that crossed application, data, and infrastructure boundaries.